



Configure FabricPool

StorageGRID

NetApp
May 03, 2022

Table of Contents

- Configure FabricPool 1
 - Configure StorageGRID for FabricPool: Overview 1
 - Attach StorageGRID as a cloud tier 3
 - Use StorageGRID information lifecycle management with FabricPool data 14
 - Create a traffic classification policy for FabricPool 16
 - Other best practices for StorageGRID and FabricPool 19

Configure FabricPool

Configure StorageGRID for FabricPool: Overview


If you use NetApp ONTAP software, you can use NetApp FabricPool to tier inactive, or cold, data to a NetApp StorageGRID object storage system.

About these instructions

Use these instructions to:

- Get an overview of configuring a StorageGRID object storage system for use with FabricPool.
- Learn how to obtain the information you provide to ONTAP when you attach StorageGRID as a FabricPool cloud tier.
- Learn the best practices for configuring the StorageGRID information lifecycle management (ILM) policy, a StorageGRID traffic classification policy, and other StorageGRID options for a FabricPool workload.

Configuration workflow



Before you start

- Decide which FabricPool volume tiering policy you will use to tier inactive ONTAP data to StorageGRID.
- Plan and install a StorageGRID system to meet your storage capacity and performance needs.
- Become familiar with StorageGRID system software, including the Grid Manager and the Tenant Manager.
- Review these additional resources, which provide details about using and configuring FabricPool:

- [TR-4598: FabricPool Best Practices in ONTAP 9.9.1](#)
- [ONTAP 9 Documentation](#)

What is FabricPool?

FabricPool is an ONTAP hybrid storage solution that uses a high-performance flash aggregate as the performance tier and an object store as the cloud tier. Data is stored either on the primary storage media or in the object datastore based on whether it is accessed frequently or not. Using FabricPool-enabled aggregates helps you reduce storage cost without compromising performance, efficiency, or protection.

No architectural changes are required, and you can continue managing your data and application environment from the central ONTAP storage system.

What is StorageGRID?

StorageGRID is a storage architecture that manages data as objects, as opposed to other storage architectures such as file or block storage. Objects are kept inside a single container (such as a bucket) and are not nested as files inside a directory inside other directories. Although object storage generally provides lower performance than file or block storage, it is significantly more scalable. StorageGRID buckets can hold petabytes of data and billions of objects.

Why use StorageGRID as a FabricPool cloud tier?

FabricPool can tier ONTAP data to a number of object store providers, including StorageGRID. Unlike public clouds that might set a maximum number of supported input/output operations per second (IOPS) at the bucket or container level, StorageGRID performance scales with the number of nodes in a system. Using StorageGRID as a FabricPool cloud tier allows you to keep your cold data in your own private cloud for highest performance and complete control over your data.

In addition, a FabricPool license is not required when you use StorageGRID as the cloud tier.

Can I use multiple ONTAP clusters with StorageGRID?

These instructions describe how to connect StorageGRID to a single ONTAP cluster. However, you might want to connect the same StorageGRID system to multiple ONTAP clusters.

The only requirement for tiering data from multiple ONTAP clusters to a single StorageGRID system is that you must use a different S3 bucket for each cluster. Based on your requirements, you can use the same high availability (HA) group, load balancer endpoint, and tenant account for all clusters, or you can configure each of these items for each cluster.

Attach StorageGRID as a cloud tier

Information needed to attach StorageGRID as a cloud tier

Before you can attach StorageGRID as a cloud tier for FabricPool, you must perform some configuration steps in StorageGRID and obtain certain values.

About this task

The following table lists the information you must provide to ONTAP when you attach StorageGRID as a cloud tier for FabricPool. The topics in this section explain how to use the StorageGRID Grid Manager and Tenant Manager to obtain the information you need.



The exact field names listed and the process you use to enter the required values in ONTAP depend on whether you are using the ONTAP CLI (storage aggregate object-store config create) or ONTAP System Manager (**Storage > Aggregates & Disks > Cloud Tier**).

For more information, refer to the following:

- [TR-4598: FabricPool Best Practices in ONTAP 9.9.1](#)
- [ONTAP 9 Documentation](#)

ONTAP field	Description
Object store name	Any unique and descriptive name. For example, StorageGRID_Cloud_Tier.
Provider type	StorageGRID (ONTAP System Manager) or SGWS (ONTAP CLI).
Port	<p>The port that FabricPool will use when it connects to StorageGRID. You determine which port number to use when you define the StorageGRID load balancer endpoint.</p> <p>Create a load balancer endpoint for FabricPool</p>
Server name	<p>The fully qualified domain name (FQDN) for the StorageGRID load balancer endpoint. For example, s3.storagegrid.company.com.</p> <p>Note the following:</p> <ul style="list-style-type: none">• The domain name that you specify here must match the domain name on the CA certificate you upload for the StorageGRID load balancer endpoint.• The DNS record for this domain name must map to each IP address you will use to connect to StorageGRID. <p>Configure the DNS server for StorageGRID IP addresses</p>
Container name	<p>The name of the StorageGRID bucket you will use with this ONTAP cluster. For example, fabricpool-bucket. You can create this bucket in the Tenant Manager or, starting with ONTAP 9.10 System Manager, you can create the bucket with the FabricPool setup wizard.</p> <p>Note the following:</p> <ul style="list-style-type: none">• The bucket name cannot be changed once the configuration is created.• The bucket cannot have versioning enabled.• You must use a different bucket for each ONTAP cluster that will tier data to StorageGRID. <p>Create an S3 bucket and obtain an access key</p>

ONTAP field	Description
Access key and secret password	<p>The access key and secret access key for the StorageGRID tenant account.</p> <p>You generate these values in the Tenant Manager.</p> <p>Create an S3 bucket and obtain an access key</p>
SSL	Must be enabled.
Object store certificate	<p>The CA certificate you uploaded when you created the StorageGRID load balancer endpoint.</p> <p>Note: If an intermediate CA issued the StorageGRID certificate, you must provide the intermediate CA certificate. If the StorageGRID certificate was issued directly by the Root CA, you must provide the Root CA certificate.</p> <p>Create a load balancer endpoint for FabricPool</p>

After you finish

After you have obtained the required StorageGRID information, you can go to ONTAP to add StorageGRID as a cloud tier, add the cloud tier as an aggregate, and set volume tiering policies.

Best practices for load balancing

Before attaching StorageGRID as a FabricPool cloud tier, you must use the StorageGRID Grid Manager to configure at least one load balancer endpoint.

What is load balancing?

When data is tiered from FabricPool to a StorageGRID system, StorageGRID uses a load balancer to manage the ingest and retrieval workload. Load balancing maximizes speed and connection capacity by distributing the FabricPool workload across multiple Storage Nodes.

The StorageGRID Load Balancer service is installed on all Admin Nodes and all Gateway Nodes and provides Layer 7 load balancing. It performs Transport Layer Security (TLS) termination of client requests, inspects the requests, and establishes new secure connections to the Storage Nodes.

The Load Balancer service on each node operates independently when forwarding client traffic to the Storage Nodes. Through a weighting process, the Load Balancer service routes more requests to Storage Nodes with higher CPU availability.

Although the StorageGRID Load Balancer service is the recommended load balancing mechanism, you might want to integrate a third-party load balancer instead. For information, contact your NetApp account representative or refer to [TR-4626: StorageGRID third-party and global load balancers](#).



The separate Connection Load Balancer (CLB) service on Gateway Nodes is deprecated and no longer recommended for use with FabricPool.

Best practices for StorageGRID load balancing

As a general best practice, each site in your StorageGRID system should include two or more nodes with the Load Balancer service. For example, a site might include two Gateway Nodes or both an Admin Node and a Gateway Node. Make sure that there is adequate networking, hardware, or virtualization infrastructure for each load-balancing node, whether you are using SG100 or SG1000 services appliances, bare metal nodes, or virtual machine (VM) based nodes.

You must configure a StorageGRID load balancer endpoint to define the port that Gateway Nodes and Admin Nodes will use for incoming and outgoing FabricPool requests.

Best practices for the load balancer endpoint certificate

When creating a load balancer endpoint for use with FabricPool, you should use HTTPS as the protocol. Communicating with StorageGRID without TLS encryption is supported but not recommended.

You can then either upload a certificate that is signed by either a publicly trusted or a private certificate authority (CA), or you can generate a self-signed certificate. The certificate allows ONTAP to authenticate with StorageGRID.

As a best practice, you should use a CA server certificate to secure the connection. Certificates signed by a CA can be rotated nondisruptively.

When requesting a CA certificate for use with the load balancer endpoint, ensure that the domain name on the certificate matches the server name you enter in ONTAP for that load balancer endpoint. If possible, use a wildcard (*) to allow for virtual-host-style URLs. For example:

```
*.s3.storagegrid.company.com
```

When you add StorageGRID as a FabricPool cloud tier, you must install the same certificate to the ONTAP cluster, as well as the root and any subordinate certificate authority (CA) certificates.



StorageGRID uses server certificates for a number of purposes. If you are connecting to the Load Balancer service, you can optionally use the S3 and Swift API certificate.

To learn more about the server certificate for a load balancing endpoint:

- [Configure load balancer endpoints](#)
- [Hardening guidelines for server certificates](#)

Best practices for high availability groups

Before attaching StorageGRID as a FabricPool cloud tier, you should use the StorageGRID Grid Manager to configure a high availability (HA) group.

What is a high availability (HA) group?

To ensure that the Load Balancer service is always available to manage FabricPool data, you can group the network interfaces of multiple Admin and Gateway Nodes into a single entity, known as a high availability (HA) group. If the active node in the HA group fails, another node in the group can continue to manage the workload.

Each HA group provides highly available access to the shared services on the associated nodes. For example, an HA group that consists of interfaces only on Gateway Nodes or on both Admin Nodes and Gateway Nodes provides highly available access to the shared Load Balancer service.

To create an HA group, you perform these general steps:

1. Select network interfaces for one or more Admin Nodes or Gateway Nodes. You can select the Grid Network interface (eth0), Client Network interface (eth2), or a VLAN interface.



If you plan to use a VLAN interface to segregate FabricPool traffic, a network administrator must first configure a trunk interface and the corresponding VLAN. Each VLAN is identified by a numeric ID or tag. For example, your network might use VLAN 100 for FabricPool traffic.

2. Assign one or more virtual IP (VIP) addresses to the group. Clients applications, such as FabricPool, can use any of these VIP addresses to connect to StorageGRID.
3. Specify one interface to be the Primary interface and determine the priority order for any Backup interfaces. The Primary interface is the active interface unless a failure occurs.

If the HA group includes more than one interface and the Primary interface fails, the VIP addresses move to the first backup interface in the priority order. If that interface fails, the VIP addresses move to the next backup interface, and so on. This failover process generally takes only a few seconds and is fast enough that client applications should experience little impact and can rely on normal retry behaviors to continue operation.

When the failure is resolved and a higher priority interface becomes available again, the VIP addresses are automatically moved to the highest priority interface that is available.

Best practices for high availability (HA) groups

The best practices for creating a StorageGRID HA group for FabricPool depend on the workload, as follows:

- If you plan to use FabricPool with primary workload data, you must create a HA group that includes at least two load-balancing nodes to prevent data retrieval interruption.
- If you plan to use the FabricPool snapshot-only volume tiering policy or non-primary local performance tiers (for example, disaster recovery locations or NetApp SnapMirror® destinations), you can configure an HA group with only one node.

These instructions describe setting up an HA group for Active-Backup HA (one node is active and one node is backup). However, you might prefer to use DNS Round Robin or Active-Active HA. To learn the benefits of these other HA configurations, see [Configuration options for HA groups](#).

Configure the DNS server for StorageGRID IP addresses

After configuring high availability groups and load balancer endpoints, you must ensure that the domain name system (DNS) for the ONTAP system includes a record to associate the StorageGRID server name (fully qualified domain name) to the IP address that FabricPool will use to make connections.

The IP address you enter in the DNS record depends on whether you are using an HA group of load-balancing nodes:

- If you have configured a HA group, FabricPool will connect to the virtual IP addresses of that HA group.

- If you are not using a HA group, FabricPool can connect to the StorageGRID Load Balancer service using the IP address of any Gateway Node or Admin Node.

You must also ensure that the DNS record references all required endpoint domain names, including any wildcard names.

Create a high availability (HA) group for FabricPool

When configuring StorageGRID for use with FabricPool, you can optionally create one or more high availability (HA) groups. An HA group consists of one or more network interfaces on Admin Nodes, Gateway Nodes, or both.

What you'll need

- You are signed in to the Grid Manager using a [supported web browser](#).
- You have the Root Access permission.
- If you plan to use a VLAN, you have created the VLAN interface. See [Configure VLAN interfaces](#).

About this task

Each HA group uses virtual IP addresses (VIPs) to provide highly available access to the shared services on the associated nodes.

For details about this task, see [Manage high availability groups](#).

Steps

1. Select **CONFIGURATION > Network > High availability groups**.
2. Select **Create**.
3. Enter a unique name and optionally a description.
4. Select one or more interfaces to add to this HA group.

Use the column headers to sort the rows, or enter a search term to locate interfaces more quickly.

5. Determine the Primary interface and any backup interfaces for this HA group.

Drag and drop rows to change the values in the **Priority order** column.

The first interface in the list is the Primary interface. The Primary interface is the active interface unless a failure occurs.

If the HA group includes more than one interface and the active interface fails, the VIP addresses move to the first backup interface in the priority order. If that interface fails, the VIP addresses move to the next backup interface, and so on. When failures are resolved, the VIP addresses move back to highest priority interface available.

6. Specify the VIP subnet in CIDR notation—an IPv4 address followed by a slash and the subnet length (0-32).

The network address must not have any host bits set. For example, 192.16.0.0/22.

7. Optionally, if the ONTAP IP addresses used to access StorageGRID are not on the same subnet as the StorageGRID VIP addresses, enter the StorageGRID VIP local gateway IP address. The local gateway IP address must be within the VIP subnet.

8. Enter one or more virtual IP addresses for the HA group. You can add up to 10 IP addresses. All VIPs must be within the VIP subnet.

You must provide at least one IPv4 address. Optionally, you can specify additional IPv4 and IPv6 addresses.

9. Select **Create HA group** and then select **Finish**.

Create a load balancer endpoint for FabricPool

When configuring StorageGRID for use with FabricPool, you must configure a load balancer endpoint and upload the load balancer endpoint certificate, which is used to secure the connection between ONTAP and StorageGRID.

What you'll need

- You are signed in to the Grid Manager using a [supported web browser](#).
- You have the Root access permission.
- You have the following files:
 - Server Certificate: The custom server certificate file.
 - Server Certificate Private Key: The custom server certificate private key file.
 - CA Bundle: A single optional file containing the certificates from each intermediate issuing certificate authority (CA). The file should contain each of the PEM-encoded CA certificate files, concatenated in certificate chain order.

About this task

For details about this task, see [Configure load balancer endpoints](#).

Steps

1. Select **CONFIGURATION > Network > Load balancer endpoints**.
2. Select **Create**.

×

Create a load balancer endpoint

1 Enter endpoint details

2 Select binding mode

3 Attach certificate

Endpoint details

Name ?

Port ?

Enter an unused port or accept the suggested port.

10443

Client type ?

Select the type of client application that will use this endpoint.

☒ S3
 ☐ Swift

Network protocol ?

Select the network protocol clients will use with this endpoint. If you select HTTPS, attach the security certificate before saving the endpoint.

☐ HTTPS (recommended)
 ☒ HTTP

Cancel

Continue

3. Enter endpoint details.

Field	Description
Name	A descriptive name for the endpoint
Port	<p>The StorageGRID port you want to use for load balancing. This field defaults to 10433, but you can enter any unused external port. If you enter 80 or 443, the endpoint is configured only on Gateway Nodes, since these ports are reserved on Admin Nodes.</p> <p>Note: Ports used by other grid services are not permitted. See the Network port reference.</p> <p>You must provide this same port number to ONTAP when you attach StorageGRID as a FabricPool cloud tier.</p>
Client type	Select S3 .
Network protocol	<p>Select HTTPS.</p> <p>Note: Using HTTP is supported but not recommended.</p>

4. Select **Continue**.
5. Specify the binding mode.

Use the **Global** setting (recommended) or restrict the accessibility of this endpoint to one of the following:

- Specific network interfaces of specific nodes.
- Specific high availability (HA) virtual IP addresses (VIPs). Use this selection only if you require much higher levels of isolation of workloads.

6. Select **Continue**.
7. Select **Upload Certificate** (recommended) and then browse to your server certificate, certificate private key, and optional CA bundle.
8. Select **Create**.



Changes to an endpoint certificate can take up to 15 minutes to be applied to all nodes.

Create a tenant account for FabricPool

You must create a tenant account in the Grid Manager for FabricPool use.

What you'll need

- You are signed in to the Grid Manager using a [supported web browser](#).
- You have specific access permissions.

About this task

Tenant accounts allow client applications to store and retrieve objects on StorageGRID. Each tenant account has its own account ID, authorized groups and users, buckets, and objects.

You can use the same tenant account for multiple ONTAP clusters. Or, you can create a dedicated tenant account for each ONTAP cluster as required.



These instructions assume that you have configured single sign-on (SSO) for the Grid Manager. If SSO is not enabled, use [these instructions to create a tenant account](#) instead.

Steps

1. Select **TENANTS**.
2. Select **Create**.
3. Enter a display name and a description.
4. Select **S3**.
5. Leave the **Storage quota** field blank.
6. Select **Allow platform services** to enable the use of platform services.

If platform services are enabled, a tenant can use features, such as CloudMirror replication, that access external services.

7. Do not select **Use own identity source**.
8. Do not select **Allow S3 Select**.

9. Select an existing federated group from the Grid Manager to have the initial Root Access permission for the tenant.
10. Select **Create tenant**.

Create an S3 bucket and obtain an access key

Before using StorageGRID with a FabricPool workload, you must create an S3 bucket for your FabricPool data. You also need to obtain an access key and secret access key for the tenant account you will use for FabricPool.

What you'll need

- You have created a tenant account for FabricPool use.

About this task

These instructions describe how to use the StorageGRID Tenant Manager to create a bucket and obtain access keys. You can also perform these tasks using the Tenant Management API or the StorageGRID S3 REST API. Or, if you are using ONTAP 9.10, you can create the bucket using the FabricPool setup wizard instead.

To learn more:

- [Use a tenant account](#)
- [Use S3](#)

Steps

1. Sign in to the Tenant Manager.

You can do either of the following:

- From the Tenant Accounts page in the Grid Manager, select the **Sign in** link for the tenant, and enter your credentials.
- Enter the URL for the tenant account in a web browser, and enter your credentials.

2. Create an S3 bucket for FabricPool data.

You must create a unique bucket for each ONTAP cluster you plan to use.

- a. Select **STORAGE (S3) > Buckets**.
- b. Select **Create bucket**.
- c. Enter the name of the StorageGRID bucket you will use with FabricPool. For example, `fabricpool-bucket`.



You cannot change the bucket name after creating the bucket.

Bucket names must comply with these rules:

- Must be unique across each StorageGRID system (not just unique within the tenant account).
- Must be DNS compliant.
- Must contain at least 3 and no more than 63 characters.

- Can be a series of one or more labels, with adjacent labels separated by a period. Each label must start and end with a lowercase letter or a number and can only use lowercase letters, numbers, and hyphens.
- Must not look like a text-formatted IP address.
- Should not use periods in virtual hosted style requests. Periods will cause problems with server wildcard certificate verification.

d. Select the region for this bucket.

By default, all buckets are created in the `us-east-1` region.

Create bucket
×

Enter bucket details

Enter the bucket's name and select the bucket's region.

Bucket name ?

fabricpool-bucket

Region ?

us-east-1 ▼

Cancel
Create bucket

e. Select **Create bucket**.



For FabricPool buckets, the recommended bucket consistency level is “read-after-new-write.” The “available” consistency level is not recommended. See [Change the consistency level](#).

3. Create an access key and a secret access key.

- Select **STORAGE (S3) > My access keys**.
- Select **Create key**.
- Select **Create access key**.
- Copy the access key ID and the secret access key to a safe location, or select **Download .csv** to save a spreadsheet file containing the access key ID and secret access key.

You will enter these values in ONTAP when you configure StorageGRID as a FabricPool cloud tier.



If you create a new access key and secret access key in the future, remember to update the corresponding values in ONTAP immediately to ensure that ONTAP can store and retrieve data in StorageGRID without interruption.

Use StorageGRID information lifecycle management with FabricPool data

If you are using FabricPool to tier data to StorageGRID, you must understand the requirements for creating StorageGRID information lifecycle management (ILM) rules and an ILM policy to manage FabricPool data. You must ensure the ILM rules that apply to FabricPool data are not disruptive.



FabricPool has no knowledge of StorageGRID ILM rules or policies. Data loss can occur if the StorageGRID ILM policy is misconfigured. See [Manage objects with ILM](#) for detailed ILM instructions.

Review these guidelines to ensure that your ILM rules and ILM policy are suitable for FabricPool data and your business requirements. If you are already using StorageGRID ILM, you might need to update your active ILM policy to meet these guidelines.

- You can use any combination of replication and erasure-coding rules to protect cloud tier data.

The recommended best practice is to use 2+1 erasure coding within a site for cost-efficient data protection. Erasure coding uses more CPU, but offers significantly less storage capacity, than replication. The 4+1 and 6+1 schemes use less capacity than the 2+1 scheme. However, the 4+1 and 6+1 schemes are less flexible if you need to add Storage Nodes during grid expansion. For details, see [Add storage capacity for erasure-coded objects](#).

- Each rule applied to FabricPool data must either use erasure coding or it must create at least two replicated copies.



An ILM rule that creates only one replicated copy for any time period puts data at risk of permanent loss. If only one replicated copy of an object exists, that object is lost if a Storage Node fails or has a significant error. You also temporarily lose access to the object during maintenance procedures such as upgrades.

- Do not use an ILM rule that will expire or delete FabricPool cloud tier data. Set the retention period in each ILM rule to "forever" to ensure that FabricPool objects are not deleted by StorageGRID ILM.
- Do not create rules that will move FabricPool cloud tier data out of the bucket to another location. You cannot use ILM rules to archive FabricPool data to tape using an Archive Node or use a Cloud Storage Pool to move FabricPool data to another object store.



Using Cloud Storage Pools with FabricPool is not supported because of the added latency to retrieve an object from the Cloud Storage Pool target.

- Starting with ONTAP 9.8, you can optionally create object tags to help classify and sort tiered data for easier management. For example, you can set tags only on FabricPool volumes attached to StorageGRID. Then, when you create ILM rules in StorageGRID, you can use the Object Tag advanced filter to select and place this data.

Example ILM policy for FabricPool data

Use this simple example policy as a starting point for your own ILM rules and policy.

This example assumes you are designing the ILM rules and an ILM policy for a StorageGRID system that has

four Storage Nodes at a single data center in Denver, Colorado. The FabricPool data in this example uses a bucket named `fabricpool-bucket`.



The following ILM rules and policy are only examples. There are many ways to configure ILM rules. Before activating a new policy, simulate the proposed policy to confirm it will work as intended to protect content from loss. To learn more, see [Manage objects with ILM](#).

Steps

- 1. Create a storage pool named **DEN**. Select the Denver site.
- 2. Create an Erasure Coding profile named **2 plus 1**. Select the 2+1 erasure-coding scheme and the **DEN** storage pool.
- 3. Create an ILM rule that applies only to the data in `fabricpool-bucket`. This example rule creates erasure-coded copies.

Rule definition	Example value
Rule Name	2 plus 1 erasure coding for FabricPool data
Bucket Name	<code>fabricpool-bucket</code> You could also filter on the FabricPool tenant account.
Advanced Filtering	Object Size (MB) greater than 0.2 MB. Note: FabricPool only writes 4 MB objects, but you must add an Object Size filter because this rule uses erasure coding.
Reference Time	Ingest Time
Placement	From day 0 store forever
Type	Erasure coded
Location	DEN (2 plus 1)
Ingest Behavior	Balanced

- 4. Create an ILM rule that will create two replicated copies of any objects not matched by the first rule. Do not select a basic filter (tenant account or bucket name) or any advanced filters.

Rule definition	Example value
Rule Name	Two replicated copies
Bucket Name	<i>none</i>
Advanced Filtering	<i>none</i>

Rule definition	Example value
Reference Time	Ingest Time
Placement	From day 0 store forever
Type	Replicated
Location	DEN
Copies	2
Ingest Behavior	Balanced

5. Create a proposed ILM policy and select the two rules. Because the replication rule does not use any filters, it can be the default (last) rule for the policy.
6. Ingest test objects into the grid.
7. Simulate the policy with the test objects to verify the behavior.
8. Activate the policy.

When this policy is activated, StorageGRID places object data as follows:

- The data tiered from FabricPool in `fabricpool-bucket` will be erasure coded using the 2+1 erasure-coding scheme. Two data fragments and one parity fragment will be placed on three different Storage Nodes.
- All objects in all other buckets will be replicated. Two copies will be created and placed on two different Storage Nodes.
- The erasure-coded and replicated copies will be maintained in StorageGRID until they are deleted by the S3 client. StorageGRID ILM will never delete these items.

Create a traffic classification policy for FabricPool

You can optionally design a StorageGRID traffic classification policy to optimize quality of service for the FabricPool workload.

What you'll need

- You are signed in to the Grid Manager using a [supported web browser](#).
- You have the Root Access permission.

About this task

The best practices for creating a traffic classification policy for FabricPool depend on the workload, as follows:

- If you plan to tier FabricPool primary workload data to StorageGRID, you should ensure that the FabricPool workload has the majority of bandwidth. You can create a traffic classification policy to limit all other workloads.



In general, FabricPool read operations are more important to prioritize than write operations.

For example, if other S3 clients use this StorageGRID system, you should create a traffic classification policy. You can limit network traffic for the other buckets, tenants, IP subnets, or load balancer endpoints.

- As a general rule, you should not impose quality of service limits on any FabricPool workload; you should only limit the other workloads.
- The limits placed on other workloads should account for the behavior of those workloads. The limits imposed will also vary based on the sizing and capabilities of your grid and what the expected amount of utilization is.

To learn more: [Manage traffic classification policies](#)

Steps

1. Select **CONFIGURATION > Network > Traffic classification**.
2. Enter a name and a description.
3. In the Matching Rules section, create at least one rule.
 - a. Select **Create**.
 - b. Select **Endpoint**, and select the load balancer endpoint you created for FabricPool.

You can also select the FabricPool tenant account or bucket.

- c. If you want this traffic policy to limit traffic for the other endpoints, select **Inverse Match**.
4. Optionally, create one or more limits.



Even if no limits are set for a traffic classification policy, metrics are collected so you can understand traffic trends.

- a. Select **Create**.
- b. Select the type of traffic you want to limit and the limit to apply.

This example FabricPool traffic classification policy shows the types of network traffic you can limit and the types of values you can select. The limits for an actual policy would be based on your specific requirements.

Policy

Name

Description (optional)

Matching Rules

Traffic that matches any rule is included in the policy.

+ Create
Edit
Remove

Type	Inverse Match	Match Value
<input checked="" type="radio"/> Endpoint	<input checked="" type="checkbox"/>	FabricPool (https 10443)

Displaying 1 matching rule.

Limits (Optional)

+ Create
Edit
Remove

Type	Value	Units
<input type="radio"/> Concurrent Read Requests	50	Concurrent Requests
<input type="radio"/> Concurrent Read Requests	15	Concurrent Requests
<input type="radio"/> Read Request Rate	100	Requests/Second
<input type="radio"/> Write Request Rate	25	Requests/Second
<input type="radio"/> Per-Request Bandwidth In	2000000	Bytes/Second
<input checked="" type="radio"/> Per-Request Bandwidth Out	10000000	Bytes/Second

- After creating the traffic classification policy, select the policy and then select **Metrics** to determine if the policy is limiting traffic as expected.

Traffic Classification Policies

Traffic classification policies can be used to identify network traffic for metrics reporting and optional traffic limiting.

+ Create Edit Remove Metrics		
Name	Description	ID
<input checked="" type="radio"/> FabricPool	Limit traffic other than FabricPool	587f53b2-7cf2-44b9-af5c-694ebbd4a2c5
Displaying 1 traffic classification policy.		

Other best practices for StorageGRID and FabricPool

When configuring a StorageGRID system for use with FabricPool, you should avoid setting global options that might affect how your data is saved.

Object encryption

When configuring StorageGRID, you can optionally enable the global **Stored Object Encryption** setting if data encryption is required for other StorageGRID clients (**CONFIGURATION > System > Grid options**). The data that is tiered from FabricPool to StorageGRID is already encrypted, so enabling the StorageGRID setting is not required. Client-side encryption keys are owned by ONTAP.

Object compression

When configuring StorageGRID, do not enable the global **Compress Stored Objects** setting (**CONFIGURATION > System > Grid options**). The data that is tiered from FabricPool to StorageGRID is already compressed. Enabling **Compress Stored Objects** will not further reduce an object's size.

Consistency level

For FabricPool buckets, the recommended bucket consistency level is “read-after-new-write.” The “available” consistency level is not recommended. See [Change the consistency level](#).

FabricPool tiering

If the StorageGRID node uses storage assigned from a NetApp AFF system, confirm that the volume does not have a FabricPool tiering policy enabled. For example, if a StorageGRID node is running on a VMware host, ensure the volume backing the datastore for the StorageGRID node does not have a FabricPool tiering policy enabled. Disabling FabricPool tiering for volumes used with StorageGRID nodes simplifies troubleshooting and storage operations.



Never use FabricPool to tier any data related to StorageGRID back to StorageGRID itself. Tiering StorageGRID data back to StorageGRID increases troubleshooting and operational complexity.

Copyright Information

Copyright © 2022 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.